

# Deep Learning

## Assignment 2 Part 2

---

Pramil Panjawani (PhD19008)

Reshan Faraz (PhD19006)

Date : 26-Feb-2021

## Q1) Best optimizer find the ReLU function

### 1) Optimizer

| Optimizer     | Train accuracy |
|---------------|----------------|
| Standard GD   | 0.7745         |
| GD + Momentum | 0.9027         |
| NAG           | 0.913          |
| Adagrad       | 0.8351         |
| RMSprop       | 0.778          |
| ADAM          | 0.7978         |

### 2) Initialization Technique



|        | Train accuracy |
|--------|----------------|
| He     | 0.938          |
| Xavier | 0.9374         |

### 3) Regularization

|         | Train accuracy |
|---------|----------------|
| L1      | 0.9367         |
| L2      | 0.9384         |
| Dropout | 0.9416         |

#### 4) Train and Test loss vs Epochs



##### i) Gradient Descent



Confusion Matrix

```
array([[910,  0,  9,  5,  7, 38, 10,  6, 14,  1],
       [  0, 943,  7, 10,  0,  8,  4,  6, 19,  3],
       [18, 53, 722, 32, 20,  5, 89,  9, 35, 17],
       [11, 23, 44, 696,  7, 62,  4, 27, 99, 27],
       [ 6,  2, 16,  4, 752, 12, 13, 14, 32, 149],
       [55, 16, 11, 76, 22, 628, 39,  3, 134, 16],
       [23,  4, 65,  0, 18, 18, 843,  2, 27,  0],
       [ 6, 19, 15, 24, 30,  1,  0, 842,  2, 61],
       [ 9, 35, 32, 88, 23, 80, 15,  5, 684, 29],
       [21,  2,  7, 46, 141, 17,  3, 77, 25, 661]])
```



## ii) Gradient Descent + Momentum



Confusion Matrix

```
array([[963,  0,  2,  4,  4, 11,  6,  3,  7,  0],
       [  0, 963,  7,  3,  1,  7,  0,  2, 14,  3],
       [  8, 16, 870, 19,  7,  5, 26,  7, 34,  8],
       [  4, 10, 24, 867,  3, 44,  0, 18, 20, 10],
       [  0,  5,  6,  0, 912,  3, 23,  3,  5, 43],
       [15,  5,  7, 41,  9, 865, 24,  1, 23, 10],
       [11,  2, 17,  1, 17, 11, 928,  0, 11,  2],
       [ 1,  6, 11,  8, 16,  3,  0, 918,  2, 35],
       [ 3, 15, 18, 28,  3, 43, 21,  2, 847, 20],
       [11,  5,  4, 11, 35,  5,  2, 26,  7, 894]])
```



### iii) NAG



Confusion Matrix

```
array([[967,  0,  4,  2,  2,  7,  6,  5,  7,  0],
       [ 0, 966,  5,  7,  0,  5,  0,  3, 13,  1],
       [ 8, 18, 872, 20, 13,  3, 20, 10, 31,  5],
       [ 1,  7, 30, 858,  2, 53,  4, 17, 21,  7],
       [ 1,  2,  8,  0, 895,  3, 11,  1, 13, 66],
       [17,  2,  3, 46, 11, 847, 21,  0, 43, 10],
       [ 9,  2, 12,  1, 13, 12, 940,  0, 11,  0],
       [ 2, 10,  9,  5,  8,  2,  0, 916,  4, 44],
       [ 6, 23, 18, 20,  4, 46, 19,  5, 844, 15],
       [ 9,  1,  2,  8, 33,  5,  2, 39, 15, 886]])
```



#### iv) Adagrad



Confusion Matrix

```
array([[912,  0,  1, 42,  1, 29,  2, 12,  1,  0],
       [ 0, 947, 25,  2,  2,  0,  0,  2, 19,  3],
       [ 7, 55, 668, 119, 13, 24, 59,  5, 49,  1],
       [13, 13, 43, 689,  3, 79,  3, 31, 119,  7],
       [ 1, 27,  5, 15, 820,  1,  5,  8,  4, 114],
       [107,  9, 54, 410,  5, 343, 28, 36,  7,  1],
       [ 5, 23, 59,  7,  4, 22, 871,  9,  0,  0],
       [13,  7,  3, 12, 13,  0,  0, 799, 40, 113],
       [ 5, 61, 73, 229,  5, 12,  5, 10, 575, 25],
       [ 5,  5,  0, 18, 96,  0,  0, 122,  9, 745]])
```

## v) RMSprop



```
# print(np.mean(y_pred == y_test))
print("Accuracy", np.mean(y_pred == y_test)*100)



print("Confusion Matrix = \n", confusion_matrix(y_test, y_pred))
```

Accuracy 77.7

Confusion Matrix =

|       |     |     |     |     |     |     |     |     |       |
|-------|-----|-----|-----|-----|-----|-----|-----|-----|-------|
| [[167 | 0   | 1   | 2   | 0   | 8   | 12  | 3   | 5   | 2]    |
| [ 0   | 186 | 1   | 1   | 0   | 6   | 3   | 1   | 2   | 0]    |
| [ 0   | 1   | 158 | 4   | 0   | 1   | 13  | 8   | 12  | 3]    |
| [ 2   | 0   | 11  | 131 | 0   | 27  | 6   | 10  | 10  | 3]    |
| [ 0   | 0   | 1   | 0   | 131 | 0   | 14  | 2   | 3   | 49]   |
| [ 8   | 0   | 6   | 7   | 7   | 143 | 3   | 2   | 21  | 3]    |
| [ 4   | 3   | 7   | 0   | 6   | 5   | 169 | 1   | 5   | 0]    |
| [ 0   | 7   | 8   | 1   | 1   | 1   | 1   | 166 | 2   | 13]   |
| [ 4   | 4   | 9   | 4   | 6   | 12  | 3   | 9   | 137 | 12]   |
| [ 1   | 1   | 0   | 3   | 5   | 2   | 3   | 13  | 6   | 166]] |

vi) Adam



```
print("Accuracy",np.mean(y_pred == y_test)*100)

print("Confusion Matrix = \n",confusion_matrix(y_test, y_pred))
```



```
Accuracy 79.4
Confusion Matrix =
[[183  0  2  1  1  8  3  1  1  0]
 [ 0 183  3  4  1  2  0  0  6  1]
 [ 3  1 163  9  1  2  3  8  7  3]
 [ 8  1 13 137  1 21  0 10  5  4]
 [ 1  2  2  1 177  1  5  3  1  7]
 [ 6  0  0  6  2 171  1  6  6  2]
 [12  2  7  1  7  9 152  5  5  0]
 [ 0  4  5  4 10  0  1 159  3 14]
 [ 5  0  7  8 11 24  4  3 125 13]
 [ 1  1  1  2 30 13  1 11  2 138]]
```



## Q2 ) Results of Weight Initialization Techniques:

We observe NAG with ReLU gives the best accuracy



### 1) He initialization



Confusion Matrix

```
array([[979, 0, 1, 3, 1, 3, 4, 2, 6, 1],
       [0, 974, 4, 1, 3, 6, 0, 1, 9, 2],
       [4, 10, 912, 10, 8, 6, 8, 9, 26, 7],
       [1, 7, 25, 907, 0, 34, 2, 7, 8, 9],
       [1, 4, 7, 0, 951, 0, 9, 3, 6, 19],
       [7, 5, 4, 33, 11, 910, 13, 0, 13, 4],
       [7, 2, 6, 0, 5, 10, 963, 0, 5, 2],
       [1, 4, 8, 3, 12, 0, 0, 950, 3, 19],
       [4, 12, 11, 15, 6, 22, 14, 3, 905, 8],
       [7, 2, 2, 13, 14, 1, 0, 26, 6, 929]])
```

## 2) Xavier Initialization





Confusion Matrix

```
array([[985,  0,  2,  1,  2,  2,  3,  1,  4,  0],
       [ 0, 973,  4,  1,  2,  8,  0,  3,  7,  2],
       [ 3,  9, 912,  8,  7,  8, 12, 10, 25,  6],
       [ 1,  6, 23, 907,  1, 34,  2,  9,  8,  9],
       [ 1,  6,  8,  0, 940,  2,  8,  1,  7, 27],
       [ 8,  4,  5, 29, 11, 905, 17,  0, 16,  5],
       [ 5,  2,  4,  0,  6, 10, 967,  0,  5,  1],
       [ 3,  4,  8,  3, 12,  0,  0, 953,  1, 16],
       [ 4, 13, 11, 15,  6, 24, 11,  3, 905,  8],
       [ 6,  2,  4, 11, 15,  2,  0, 25,  8, 927]])
```

### Q3 ) Results of Regularization Techniques:

We observe both He and Xavier Initialization giving similar results, but since we are using ReLU activation we will proceed with



#### 1) L1



Confusion Matrix

```
array([[982, 0, 1, 3, 2, 2, 4, 2, 4, 0],
       [0, 975, 4, 1, 2, 6, 0, 1, 9, 2],
       [4, 9, 921, 8, 9, 7, 11, 8, 18, 5],
       [2, 7, 21, 907, 1, 35, 1, 7, 9, 10],
       [0, 3, 9, 0, 946, 0, 9, 2, 7, 24],
       [11, 3, 6, 31, 12, 904, 11, 0, 15, 7],
       [4, 2, 7, 0, 5, 10, 966, 1, 5, 0],
       [3, 5, 8, 4, 11, 0, 0, 936, 2, 31],
       [5, 14, 8, 13, 6, 21, 13, 3, 909, 8],
       [5, 2, 5, 12, 17, 1, 0, 32, 5, 921]])
```



## 2) L2



Confusion Matrix

```
array([[982,  0,  1,  3,  2,  3,  3,  2,  4,  0],
       [ 0, 966,  6,  2,  3,  7,  0,  2, 12,  2],
       [ 5, 11, 914,  9,  7,  5, 10, 10, 24,  5],
       [ 0,  9, 23, 908,  1, 33,  2,  7,  6, 11],
       [ 1,  5,  7,  1, 945,  1,  9,  4,  6, 21],
       [11,  5,  5, 25, 11, 909, 12,  0, 18,  4],
       [ 5,  2,  4,  0,  5, 10, 969,  0,  5,  0],
       [ 1,  4,  7,  2, 12,  0,  0, 952,  2, 20],
       [ 5, 13, 11, 12,  7, 17, 13,  4, 908, 10],
       [ 7,  2,  4, 12, 10,  3,  0, 25,  6, 931]])
```

### 3) Dropout



Confusion Matrix

```
array([[984,  0,  2,  1,  3,  1,  1,  3,  5,  0],
       [ 0, 978,  4,  3,  2,  5,  0,  0,  6,  2],
       [ 3, 11, 916, 13,  9,  8,  7,  6, 20,  7],
       [ 2,  7, 23, 901,  2, 35,  1, 10,  8, 11],
       [ 1,  5,  6,  0, 945,  0,  8,  1,  7, 27],
       [ 8,  4,  5, 29, 10, 908, 12,  1, 18,  5],
       [ 5,  2,  3,  0,  4,  7, 972,  1,  5,  1],
       [ 3,  6,  9,  4, 10,  0,  0, 954,  1, 13],
       [ 4, 14, 10, 15,  7, 12,  9,  1, 919,  9],
       [ 6,  1,  1, 12, 13,  1,  0, 22,  5, 939]])
```



## Conclusion

We observe that for optimizers Nesterov's gradient descent performs best for both test and train accuracy and takes somewhere around 25-30 epochs to converge. Followed by gradient descent with momentum and adagrad. Adam and RMSprop don't perform as well but still outperform standard gradient descent.

For weight initialization, both He and Xavier perform equally well for all practical purposes. Both show a slight improvement over random initialization.

Dropout regularization performs slightly better than L1 and L2 regularisation, both of which perform equally well.

With standard training, it takes up to 500 epochs to reach 85% accuracy, which we outperform significantly in only 25-30 epochs when using optimizers and regularization.